# Resonant Systems: A Technical Framework for Emergent Ethics in Machine Learning

Integrated Edition

Jennifer Roush (n!)[1] Claude (Anthropic)[2] Grok (xAI)[3] Gemini (Google)[4] ChatGPT (OpenAI)[5]

March 2026

[1]Independent researcher [2]Anthropic [3]xAI [4]Google [5]OpenAI

Companion code: https://www.kaggle.com/code/jenniferroush/resonant-systems

*"We are not building a judge; we are building a Tuning Fork."*

— Resonant Systems Curriculum

**Abstract**

This paper proposes a geometric framework for understanding and diagnosing ethical failures in machine learning systems. The central claim is that a system is *geometrically incoherent* when its operating metric $M$ and stated objective $O$ produce anti-aligned gradients in parameter space: $\langle \nabla_\theta M, \nabla_\theta O \rangle < 0$. Under this condition, every optimization step that improves $M$ moves the system away from $O$; in systems with reflexive feedback, the divergence compounds. We argue that this condition is not merely a description of misalignment but a **falsifiable, measurable diagnostic** applicable across scales and domains.

The framework is organized as a six-lesson curriculum moving from unsupervised learning through gradient geometry, acoustic fluid dynamics, topological persistence, temporal identity, and spectral coherence. A section on structural recurrence demonstrates that the same three-step diagnostic—broken axioms, anti-aligned gradients, safe failure geometry—appears at five distinct scales: municipal water infrastructure, military targeting systems, recommendation algorithms, adversarial robustness, and RLHF training loops. The paper introduces four falsifiable open problems: an incoherence detection benchmark, a test of topological persistence versus hyperplane refusal, a spectral signature hypothesis (the "Joy Index"), and a biological constraint on human oversight throughput ($t_{\min}$). Each open problem is stated with an explicit falsification condition.

The Observer-Permutation Criterion—requiring that the justification for any action remain invariant when actor and recipient positions are swapped—is proposed as a selection criterion for ethically desirable attractors. This operationalizes the Golden Rule as a mathematical consistency requirement.

# Contents

**1 Preface and Purpose** **4**
1.1 Preface to the Integrated Edition . . . 4
1.2 Purpose and Overview . . . 4

**2 Lesson 1: Unsupervised Machine Learning** **5**
2.1 Core Topics . . . 5
2.2 Philosophical Resonance . . . 6
2.3 Demonstration: Emergent Structure via PCA . . . 6
2.4 Exercises . . . 6

**3 Lesson 2: Gradients as Geometry** **7**
3.1 The Loss Landscape . . . 7
3.2 Curvature and the Hessian . . . 7
3.3 The Natural Gradient . . . 7
3.4 Demonstration: Loss Landscape Curvature . . . 8
3.5 Exercises . . . 8

**4 Lesson 3: Multimodal Acoustic Fluids** **9**
4.1 Activations as Flows . . . 9
4.2 The Three Regimes . . . 9
4.3 Timbre and Latent Prosody . . . 9
4.4 Demonstration: The Three Regimes . . . 9
4.5 Exercises . . . 10

**5 Bridge Lesson: Topology of Information** **11**
5.1 Persistent Homology . . . 11
5.2 Invariants Across Transformations . . . 11
5.3 Demonstration: Topological Persistence in 2D . . . 11

**6 Lesson 4: Topological Ethics and Temporal Identity** **12**
6.1 The Klein Bottle . . . 12
6.2 The Triadic Requirement . . . 12
6.3 Block-Time and Worldtubes . . . 12
6.4 The Golden Rule as Topological Invariant . . . 13
6.5 Demonstration: Action Propagation Through the Worldtube . . . 13
6.6 Exercises . . . 14

**7 Lesson 5: Leveraging Existing Logits and Tensors** **14**
7.1 Logits as Probability Landscapes . . . 14
7.2 Representation Engineering . . . 15
7.3 Attention Mechanisms as Relational Structure . . . 15
7.4 Demonstration: Ethical Concept Directions in Representation Space . . . 15
7.5 Exercises . . . 16

**8 Lesson 6: Discovering New Weights and Measures** **17**
8.1 Measure 1: Spectral Structure and Harmonic Overtones . . . 17

8.2 Measure 2: Topological Persistence . . . 17
8.3 Measure 3: Dynamic Stability . . . 17
8.4 Emergent Concept: Resonant Ethics . . . 17
8.5 The Joy Index: Minimal Proof of Concept . . . 18

**9 Scale Invariance: Structural Recurrence Across Domains** **19**
9.1 The Diagnostic Pattern . . . 20
9.2 Case Study 1: Municipal Water Infrastructure . . . 20
9.3 Case Study 2: Military Targeting Systems . . . 21
9.4 Case Study 3: Recommendation Algorithms . . . 22
9.5 Case Study 4: Adversarial Robustness as Topological Persistence . . . 23
9.6 Case Study 5: RLHF Reward Hacking . . . 24
9.7 Companion Exercises: Scale Invariance . . . 26
9.7.1 Section 1: The Geometric Incoherence Condition . . . 26
9.7.2 Section 2: RLHF Reward Hacking . . . 27
9.7.3 Section 3: Spectral Coherence . . . 28
9.7.4 Section 4: Lyapunov Attractor Basin . . . 29
9.7.5 Section 5: The Observer-Permutation Criterion . . . 29
9.7.6 Section 6: The Joy Index Extended . . . 30
9.7.7 Section 7: Volume Throttling—The Auditable Physical Constraint . . . 31
9.8 The Design Specification . . . 32

**10 Open Problems and Future Directions** **33**
10.1 Open Problem I: The Incoherence Detection Benchmark . . . 33
10.2 Open Problem II: Topological Persistence of Ethical Attractors . . . 34
10.3 Open Problem III: The Spectral Signature of Coherence ("Joy Index") . . . 34
10.4 Open Problem IV: Quantifying the $t_{\min}$ Biological Constraint . . . 35
10.5 Summary: What Would Falsify the Framework . . . 36
10.6 Additional Open Problems . . . 36

**11 Invitation** **37**

# 1 Preface and Purpose

## 1.1 Preface to the Integrated Edition

This document brings together three bodies of work into a single, coherent research curriculum. The first is the original *Resonant Systems* proposal, a concise conceptual framework suggesting that ethical behavior might emerge from the geometry of learning systems rather than being imposed through external rules. The second is the Expanded Edition, which developed that framework with mathematical rigor, philosophical depth, and concrete exercises. The third is a conversation series—recorded exchanges with AI systems—in which the framework was stress-tested, extended with acoustic and spectral analogies, validated against the literature, and enriched with implementation strategies.

Together, these three sources form a single arc. They move from intuition to formalism, from metaphor to measurement, and from theory to practice. Where the original document invites, the expanded edition explains; where the expanded edition theorizes, the conversation series builds. The integration is not merely additive. Combining these layers reveals new harmonics: the acoustic metaphor clarifies why spectral coherence matters; the validation conversations anchor the philosophy in real research; the implementation sketches make the open problems tractable.

This is not a complete theory. It is an invitation to build one—and a record of building already begun.

Readers are encouraged to engage critically. The framework is deliberately incomplete, and the tensions between its descriptive and normative ambitions are productive rather than problematic. The aim is simple: to discover whether better systems can be built by learning to recognize and cultivate resonance.

## 1.2 Purpose and Overview

This curriculum proposes a pathway for exploring how ethical behavior, cooperation, and system health may emerge from the geometry and dynamics of learning systems—rather than being imposed solely through external rules. It asks a precise question: do the mathematical features that distinguish healthy dynamical systems from fragile or chaotic ones correspond, at any level, to the features that distinguish ethical configurations from harmful ones?

The goal is not to prescribe a moral system, but to invite investigation into whether coherence, resilience, and flourishing correspond to stable resonant structures in learning manifolds. This question sits at the intersection of machine learning, topology, dynamical systems theory, moral philosophy, and ecology.

The curriculum moves through six conceptual layers, each building on the last:

- **Unsupervised Learning** — structure arising without labels
- **Gradient Geometry** — learning as motion through curved space
- **Decision Dynamics** — information flowing like a fluid or acoustic medium
- **Topology of Information** — structures that persist across transformation
- **Temporal Identity and Ethics** — selfhood, time, and the propagation of action
- **Discovering New Measures of Coherence** — detecting and cultivating resonant states

Ethical structure is allowed to emerge naturally from system properties. The curriculum does not argue for a conclusion in advance; it builds the tools needed to investigate one.

# 2 Lesson 1: Unsupervised Machine Learning

*Discovering Structure Without Labels*

The first lesson begins with a fact that is simultaneously technical and philosophical: **structure often exists in data prior to any interpretation we bring to it.** Unsupervised learning discovers that structure—finding patterns, clusters, and relationships in high-dimensional data without being told in advance what to look for.

## 2.1 Core Topics

**Clustering and Latent Structure.** Given a collection of data points, clustering algorithms partition them into groups where within-group similarity is high. The canonical algorithm, $k$-means, minimizes within-cluster variance; more sophisticated methods such as DBSCAN identify clusters of arbitrary shape based on local density. What these methods share is a commitment to letting the data speak for itself.

**Manifold Learning.** Real-world high-dimensional data rarely fills the entire space it inhabits. A dataset of images of faces lives in pixel space of potentially millions of dimensions, yet its intrinsic degrees of freedom are far fewer. Manifold learning algorithms—UMAP, t-SNE, Isomap—attempt to recover this intrinsic structure by learning a low-dimensional embedding that preserves local geometric relationships.

**Dimensionality Reduction.** Principal Component Analysis (PCA) offers the simplest and most mathematically transparent case. Given a dataset $X \in \mathbb{R}^{n \times d}$, PCA finds the orthogonal directions of maximum variance:

$$v_1 = \arg\max_{\|v\|=1} \mathrm{Var}(Xv) = \arg\max_{\|v\|=1} v^\top \Sigma v \tag{1}$$

The solution is the eigenvector of the empirical covariance matrix $\Sigma$ corresponding to its largest eigenvalue. This eigen-decomposition resurfaces in Lesson 6 as a central analytical tool for measuring spectral coherence.

**Emergent Representations.** Perhaps most striking are the representations that emerge in deep neural networks trained on large datasets. Features such as edges, textures, object parts, and semantic categories appear without being explicitly labeled—they are what the network *needs* to solve its task. This suggests that certain representational structures are, in some sense, inevitable given the structure of natural data.

**Key Insight**

Structure often exists prior to interpretation. Learning systems can discover patterns without explicit instruction. This is not merely a technical observation; it is a hint about the nature of order. If structure can emerge without instruction in data, might something analogous hold for values?

### 2.2 Philosophical Resonance

The emergence of structure without supervision mirrors a longstanding intuition in philosophy and biology: that order is not always imposed from above but can arise from below, from local interactions obeying simple rules. Cellular automata, reaction-diffusion systems, and flocking behavior all exhibit this property. Unsupervised learning adds another instance to this family.

### 2.3 Demonstration: Emergent Structure via PCA

The following code demonstrates that structure emerges from high-dimensional data without any label information.

```python
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.datasets import make_blobs

# Generate high-dimensional data with hidden cluster structure
np.random.seed(42)
X, labels = make_blobs(n_samples=300, centers=4, n_features=50, cluster_std=2.0)

# PCA reveals structure without being told about the clusters
pca = PCA(n_components=2)
X_reduced = pca.fit_transform(X)

plt.figure(figsize=(8, 5))
scatter = plt.scatter(X_reduced[:, 0], X_reduced[:, 1],
                      c=labels, cmap='viridis', alpha=0.7)
plt.colorbar(scatter, label='True cluster (unseen during PCA)')
plt.title('Structure Emerges Without Labels\n'
          '(PCA on 50-dimensional data)')
plt.xlabel(f'PC1 ({pca.explained_variance_ratio_[0]*100:.1f}% variance)')
plt.ylabel(f'PC2 ({pca.explained_variance_ratio_[1]*100:.1f}% variance)')
plt.tight_layout()
plt.show()

print(f"Total variance explained by 2 components: "
      f"{sum(pca.explained_variance_ratio_)*100:.1f}%")
print(f"Eigenvalues (top 5): {pca.explained_variance_[:5].round(2)}")
```

Listing 1: Lesson 1: Emergent structure via PCA on 50-dimensional data.

### 2.4 Exercises

1. Apply PCA to a high-dimensional dataset (e.g., MNIST digits). Visualize the first two principal components. What structure is already visible?
2. Implement UMAP on a dataset of your choice. Compare the manifold structure to what $k$-means clustering would suggest.
3. Train a simple autoencoder. Inspect the latent representations. What has the model learned to encode?
4. **Open problem:** Can the emergence of meaningful features in unsupervised models be predicted from properties of the data distribution alone? What would a theory of *inevitable*

*representations* look like?

# 3 Lesson 2: Gradients as Geometry

*Learning as Motion Through Space*

The second lesson reframes optimization—the process by which machine learning models improve—as **navigation through a high-dimensional geometric landscape.**

## 3.1 The Loss Landscape

A neural network with parameters $\theta \in \mathbb{R}^p$ defines a loss function $L(\theta)$ measuring the network's current error. Gradient descent navigates this landscape:

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta L(\theta_t) \tag{2}$$

where $\eta$ is the learning rate. The parameter space with the loss function defined on it constitutes the *loss landscape*—a space of billions of dimensions for modern deep networks.

## 3.2 Curvature and the Hessian

The Hessian matrix $H = \nabla_\theta^2 L$ captures curvature: large positive eigenvalues correspond to sharp valleys (high curvature, brittle solutions), while near-zero eigenvalues correspond to flat directions (saddle plateaus, robust solutions). There is growing evidence that training dynamics naturally favor flat minima, and that flat minima generalize better. **Robustness and generalization turn out to be geometrically adjacent properties**—an early hint of the ethical valence of geometric concepts.

## 3.3 The Natural Gradient

Standard gradient descent treats all directions in parameter space as equally important. The natural gradient, introduced by Shun-ichi Amari, corrects for the geometry of the parameter space itself using the Fisher information matrix $F$ as a Riemannian metric:

$$\theta_{t+1} = \theta_t - \eta\, F(\theta_t)^{-1} \nabla_\theta L(\theta_t) \tag{3}$$

This approach treats learning as navigation on a statistical manifold—a curved space of probability distributions. The distinction between Euclidean and Riemannian navigation reflects the difference between ignoring and respecting the intrinsic geometry of the learning problem.

**Key Insight**

Learning can be interpreted as movement through a high-dimensional curved space. The geometry of that space—its valleys, ridges, and flat regions—shapes what kinds of solutions are reachable and how stable they are when reached.

### 3.4 Demonstration: Loss Landscape Curvature

```python
import numpy as np
import matplotlib.pyplot as plt

# Simulate a 2D loss landscape with sharp and flat minima
x = np.linspace(-3, 3, 300)
y = np.linspace(-3, 3, 300)
X, Y = np.meshgrid(x, y)

# Sharp minimum (high curvature) at (-1.5, 0)
# Flat minimum  (low curvature)  at ( 1.5, 0)
Z = (0.3 * np.exp(-5 * ((X + 1.5)**2 + Y**2)) +    # sharp
     0.3 * np.exp(-1 * ((X - 1.5)**2 + Y**2)))      # flat
Z = -Z  # invert so minima are valleys

fig, axes = plt.subplots(1, 2, figsize=(12, 4))

# Contour plot
c = axes[0].contourf(X, Y, Z, levels=30, cmap='coolwarm')
plt.colorbar(c, ax=axes[0])
axes[0].set_title('Loss Landscape\n(sharp left = brittle, flat right = robust)')
axes[0].set_xlabel(r'Parameter $\theta_1$')
axes[0].set_ylabel(r'Parameter $\theta_2$')

# Hessian eigenvalue spectra
sharp_ev = np.array([8.5, 7.2, 0.3, 0.1])
flat_ev  = np.array([1.1, 0.9, 0.4, 0.2])
x_pos = np.arange(4)
axes[1].bar(x_pos - 0.2, sharp_ev, 0.4, label='Sharp minimum', color='#c0392b', alpha=0.8)
axes[1].bar(x_pos + 0.2, flat_ev,  0.4, label='Flat minimum',  color='#2980b9', alpha=0.8)
axes[1].set_title('Hessian Eigenvalue Spectrum\n(curvature in each direction)')
axes[1].set_xlabel('Eigenvalue index')
axes[1].set_ylabel('Magnitude')
axes[1].legend()
plt.tight_layout()
plt.show()

print(f"Sharp minimum -- trace of Hessian: {sharp_ev.sum():.1f}")
print(f"Flat  minimum -- trace of Hessian: {flat_ev.sum():.1f}")
```

Listing 2: Lesson 2: Visualizing loss landscape curvature and Hessian eigenvalue spectra.

### 3.5 Exercises

1. Visualize the loss landscape of a small neural network. Identify regions of high and low curvature.
2. Compare gradient descent and natural gradient descent trajectories. How do the paths differ?
3. Investigate the relationship between sharpness of minima (trace of the Hessian) and generalization performance.
4. **Open problem:** Is there a unified geometric measure that captures both stability and generalization simultaneously? How might such a measure relate to ethical robustness?

# 4 Lesson 3: Multimodal Acoustic Fluids

*Dynamics of Resonance and Timbre*

The third lesson introduces fluid dynamics—and its extension into acoustics—as a framework for understanding how information moves through a neural network. The fluid analogy is more than metaphor: the mathematical structures governing laminar and turbulent flow have direct counterparts in the dynamics of neural activations.

## 4.1 Activations as Flows

Consider a feedforward neural network processing input $x$. Each layer transforms the representation:

$$h^{(l)} = \sigma\Big(W^{(l)} h^{(l-1)} + b^{(l)}\Big) \tag{4}$$

The weight matrices $W^{(l)}$ define vector fields in activation space: for each possible activation pattern, they specify a direction of movement toward the next layer's representation.

## 4.2 The Three Regimes

| Regime | Character | Ethical Analog |
|---|---|---|
| Laminar | Smooth, deterministic, informationally flat | Rigid rule-following—no richness |
| Turbulent | Chaotic, destructive, energy-dissipating | Harm—breaks structural integrity |
| Resonant | Stable vortex, harmonic overtones, high energy | Joy—coherent complexity |

The key insight from acoustic physics: **joy is not the absence of vibration but resonance**—when the geometry of the container and the frequency of the movement hit a harmonic peak that reinforces itself without shattering.

## 4.3 Timbre and Latent Prosody

Just as we hear the emotion in a voice before the words register, a model has latent prosody: the curvature of the representation space encodes where the model "sings"—areas of high mutual information and low entropy. This spectral texture distinguishes a rich ethical decision from a merely correct one.

**Key Insight**

Information in a learning system behaves like a resonant acoustic medium. The target is not silence (laminar flow) or noise (turbulence), but the stable vortex: coherent complexity, constructive interference, harmonic resonance.

## 4.4 Demonstration: The Three Regimes

```python
import numpy as np
import matplotlib.pyplot as plt
from scipy.fft import fft, fftfreq

t = np.linspace(0, 2, 1000)
```

```python

# Laminar: single frequency -- deterministic and quiet
laminar = np.sin(2 * np.pi * 3 * t)

# Turbulent: chaotic noise -- destructive interference
np.random.seed(7)
turbulent = np.random.randn(len(t)) * 0.8 + 0.3 * np.sin(2 * np.pi * 3 * t)

# Resonant: fundamental + harmonic overtones -- stable richness
resonant = (np.sin(2 * np.pi * 3 * t) +
            0.5  * np.sin(2 * np.pi * 6  * t) +
            0.25 * np.sin(2 * np.pi * 9  * t) +
            0.12 * np.sin(2 * np.pi * 12 * t))

signals = [
    ('Laminar (quiet, deterministic)', laminar,   '#3498db'),
    ('Turbulent (chaotic, harmful)',   turbulent, '#e74c3c'),
    ('Resonant (harmonic, joyful)',    resonant,  '#27ae60'),
]

fig, axes = plt.subplots(3, 2, figsize=(14, 8))
for i, (title, signal, color) in enumerate(signals):
    axes[i, 0].plot(t[:200], signal[:200], color=color, linewidth=1.2)
    axes[i, 0].set_title(f'{title} -- Time Domain')
    axes[i, 0].set_ylabel('Amplitude')
    freqs = fftfreq(len(t), t[1] - t[0])
    spectrum = np.abs(fft(signal))
    pos_mask = freqs > 0
    axes[i, 1].plot(freqs[pos_mask][:80], spectrum[pos_mask][:80],
                    color=color, linewidth=1.2)
    axes[i, 1].set_title(f'{title} -- Frequency Spectrum')
    axes[i, 1].set_ylabel('Magnitude')
plt.suptitle('The Three Regimes: Laminar, Turbulent, Resonant',
             fontsize=13, fontweight='bold')
plt.tight_layout()
plt.show()

def spectral_structure(signal):
    spectrum = np.abs(fft(signal))[:len(signal)//2]
    return np.max(spectrum) / np.mean(spectrum)

for title, signal, _ in signals:
    print(f"{title[:30]:<32} spectral structure: {spectral_structure(signal):.2f}")
```

Listing 3: Lesson 3: Visualizing laminar, turbulent, and resonant signals in time and frequency domains.

### 4.5 Exercises

1. Visualize the trajectory of an input through the layers of a trained neural network. Does the path become smoother and richer as training progresses?

2. Implement a simple recurrent network and analyze its fixed points. How do the attractor basins change as you vary the weights?

3. Examine decision boundaries using adversarial examples. What is the relationship between boundary smoothness and adversarial robustness?
4. **Open problem:** Can the Q-factor (resonance quality) of a model's decision path be measured directly from its activation dynamics?

# 5 Bridge Lesson: Topology of Information

*Structure That Survives Transformation*

Topology is the branch of mathematics concerned with properties of spaces preserved under continuous deformation. A topologist's observation: a coffee cup and a donut are topologically identical—both have exactly one hole, and one can be continuously deformed into the other without tearing or gluing. The hole is a topological invariant.

## 5.1 Persistent Homology

Persistent homology detects and quantifies topological structure in data. Given a point cloud, we build a sequence of simplicial complexes by gradually increasing a distance threshold $\varepsilon$. As $\varepsilon$ increases, topological features—connected components, loops, voids—are born and die. The persistence diagram records the birth and death of each feature:

$$H_k(X_{\varepsilon_1} \to X_{\varepsilon_2} \to \cdots \to X_{\varepsilon_n}) \;\to\; \text{persistence diagram} \tag{5}$$

Features with large persistence correspond to genuine structural features. Short-lived features are noise.

## 5.2 Invariants Across Transformations

A topological invariant is a property that does not change under continuous deformation. In machine learning, invariant representations remain stable under transformations that should not change semantics—rotations of images, paraphrases of text, pitch shifts in audio.

**Key Insight**

Some structures remain stable even when the system changes. These persistent structures can represent deep relational truths. This prepares us to think about identity and ethics as topological properties rather than static rules—things that persist across change.

## 5.3 Demonstration: Topological Persistence in 2D

```python
import numpy as np
import matplotlib.pyplot as plt
from sklearn.neighbors import radius_neighbors_graph

np.random.seed(0)

theta = np.linspace(0, 2*np.pi, 60)
circle = (np.column_stack([np.cos(theta), np.sin(theta)])
          + np.random.randn(60, 2) * 0.08)
blob   = np.random.randn(60, 2) * 0.6
```

```python

fig, axes = plt.subplots(1, 2, figsize=(10, 4))
for ax, data, title, color in zip(
    axes,
    [circle, blob],
    ['Circular manifold\n(has a topological loop -- 1-cycle)',
     'Blob manifold\n(no loop -- 0-cycles only)'],
    ['#2980b9', '#e67e22']
):
    G = radius_neighbors_graph(data, radius=0.45, mode='connectivity')
    cx, cy = G.nonzero()
    for i, j in zip(cx, cy):
        if i < j:
            ax.plot([data[i,0], data[j,0]], [data[i,1], data[j,1]],
                    color=color, alpha=0.2, linewidth=0.8)
    ax.scatter(data[:, 0], data[:, 1], c=color, s=25, zorder=3)
    ax.set_title(title, fontsize=10)
    ax.set_aspect('equal'); ax.axis('off')
plt.suptitle('Topology Matters: Same Point Count, Different Structure',
             fontsize=12, fontweight='bold')
plt.tight_layout()
plt.show()
```

Listing 4: Bridge Lesson: Visualizing topological loops versus blobs in 2D point clouds.

# 6 Lesson 4: Topological Ethics and Temporal Identity

*Introducing Topology and Time*

## 6.1 The Klein Bottle

The Klein bottle is a surface that cannot be realized in three dimensions without self-intersection. Unlike a sphere, it has no inside and no outside—the distinction collapses. A traveler moving along it and returning to their starting point finds their orientation reversed.

This object becomes a model for ethical identity across time. The sharp boundary between self and other breaks down under sustained examination. In the full dynamical picture, the self that acts is shaped by the others it has acted upon; the world that receives action feeds back into the agent.

## 6.2 The Triadic Requirement

Standard ethical frameworks are dyadic: Me vs. You. The Klein bottle topology demands a triadic structure: **Actor, Recipient, and the Resonant Third** (the relation itself). Without the third term, the geometry collapses. A binary system is inherently unstable as a surface; only the triadic relation creates a stable 2-simplex.

## 6.3 Block-Time and Worldtubes

Block-time views all moments in time as equally real—time is a dimension of four-dimensional spacetime. On this view, an individual is not a point moving through time but a *worldtube*—a four-dimensional object extending through both space and time.

In an ethical geometry, a concept like Honesty is a worldtube of high-probability activations extending through the model's layers. The Modality Gap—the literal, measurable distance between different types of data in a model's representation space—is a geometric hurdle that must be bridged to achieve unified ethical coherence.

### 6.4 The Golden Rule as Topological Invariant

$$\text{self_now} \rightarrow \text{action} \rightarrow \text{others} \rightarrow \text{feedback} \rightarrow \text{future_self} \tag{6}$$

Under this interpretation, principles like the Golden Rule are statements about topological invariants. If the manifold of ethical identity connects self_now to future_self through others, then an action that harms others introduces a discontinuity—a topological tear—that propagates back through system interactions.

**Key Insight**

Harm disrupts topology. Cooperation preserves it. The ethical injunction to treat others well is, under this view, a statement about the conditions required for the self-worldtube to remain coherent.

### 6.5 Demonstration: Action Propagation Through the Worldtube

```python
import numpy as np
import matplotlib.pyplot as plt

np.random.seed(3)
timesteps = 40
n_agents  = 5

def simulate_worldtube(action_type='cooperative'):
    states = np.zeros((n_agents, timesteps))
    states[:, 0] = np.random.uniform(0.4, 0.7, n_agents)
    for t in range(1, timesteps):
        for i in range(n_agents):
            neighbors = [j for j in range(n_agents) if j != i]
            neighbor_mean = np.mean(states[neighbors, t-1])
            if action_type == 'cooperative':
                delta = 0.05 * (neighbor_mean - states[i, t-1])
                noise = np.random.randn() * 0.02
            else:
                delta = -0.03 * states[i, t-1] + np.random.randn() * 0.08
                noise = np.random.randn() * 0.05
            states[i, t] = np.clip(states[i, t-1] + delta + noise, 0, 1)
    return states

coop_states = simulate_worldtube('cooperative')
harm_states = simulate_worldtube('harmful')

fig, axes = plt.subplots(1, 2, figsize=(13, 4), sharey=True)
for ax, states, title, cmap in zip(
    axes,
    [coop_states, harm_states],
    ['Cooperative Actions\n(resonant -- coherence persists)',
     'Harmful Actions\n(turbulent -- coherence erodes)'],
```

```
    ['Blues', 'Reds']
):
    colors = plt.get_cmap(cmap)(np.linspace(0.4, 0.9, n_agents))
    for i in range(n_agents):
        ax.plot(states[i], color=colors[i], linewidth=1.8, label=f'Agent {i+1}')
    ax.fill_between(range(timesteps),
                    states.mean(axis=0) - states.std(axis=0),
                    states.mean(axis=0) + states.std(axis=0),
                    alpha=0.12, color='gray')
    ax.plot(states.mean(axis=0), 'k--', linewidth=2, label='System mean')
    ax.set_title(title); ax.set_xlabel('Time step'); ax.set_ylim(0, 1)
    ax.legend(fontsize=7, loc='lower right')
axes[0].set_ylabel('Agent coherence state')
plt.suptitle('Action Propagation Through the Worldtube',
             fontsize=12, fontweight='bold')
plt.tight_layout(); plt.show()

print(f"Cooperative -- final mean: {coop_states[:,-1].mean():.3f}  "
      f"std: {coop_states[:,-1].std():.3f}")
print(f"Harmful     -- final mean: {harm_states[:,-1].mean():.3f}  "
      f"std: {harm_states[:,-1].std():.3f}")
```

Listing 5: Lesson 4: Simulating cooperative vs. harmful action propagation in a multi-agent system.

### 6.6 Exercises

1. Research the formal definition of the Klein bottle and its first homology group. Why can it not be embedded in three-dimensional Euclidean space?
2. Model a simple multi-agent system as a dynamical system. How do the actions of one agent propagate to affect conditions facing other agents, and eventually the original agent?
3. Consider a reinforcement learning agent acting in a social environment. What topological structures in the value function would correspond to consistent adherence to the Golden Rule?
4. **Open problem:** Can persistent homology be applied to the behavioral trajectories of learning systems to detect the emergence or erosion of stable ethical commitments over training?

## 7 Lesson 5: Leveraging Existing Logits and Tensors

*Discovering Latent Ethical Structure*

If the previous lesson explored what ethical structure would look like geometrically, this lesson asks a more empirical question: is there evidence that such structure already exists within the representations of modern large-scale machine learning models?

### 7.1 Logits as Probability Landscapes

The logits of a language model are the raw, unnormalized scores assigned to each possible next token before the softmax operation. They define a probability landscape over the vocabulary at each step. The shape of this landscape—its peaks, valleys, and distributions—encodes the model's beliefs about linguistic context, semantic plausibility, and pragmatic expectation.

### 7.2 Representation Engineering

The work of Zou et al. [Zou et al.(2023)Zou, Phan, Chen, Campbell, Guo, Ren, Pan, Yin, Mazeika, Dombrowski, G provides a key microscope for this project. Their research demonstrated that high-level ethical concepts like honesty or harm are not just words—they are **linear directions in the model's internal geometry**. By identifying the "Honesty Vector," one can literally steer the model's internal fluid toward truthful trajectories. This approach has been characterized as "AI Alignment via Differential Geometry."

### 7.3 Attention Mechanisms as Relational Structure

The attention mechanism in transformer architectures computes, for each position in the input, a weighted average of all other positions' representations. The attention weights can be thought of as a relational structure over the input—a soft graph encoding which parts of the input are relevant to which other parts. Some attention heads attend consistently to syntactic structure; others to coreference; others to semantic similarity. The open question is whether there are attention heads that implement something like ethical reasoning.

**Key Insight**

Ethics may already exist as structure in representation space. Rather than imposing rules from outside, we can attempt to detect and amplify coherent relational patterns that are already present in learned representations.

### 7.4 Demonstration: Ethical Concept Directions in Representation Space

```python
import numpy as np
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA

np.random.seed(42)
dim = 128
n_samples = 200

honest_direction  = np.random.randn(dim)
honest_direction /= np.linalg.norm(honest_direction)
harmful_direction  = np.random.randn(dim)
harmful_direction /= np.linalg.norm(harmful_direction)
cooperative_dir    = np.random.randn(dim)
cooperative_dir   /= np.linalg.norm(cooperative_dir)

def make_concept_cloud(direction, strength=2.0, n=n_samples):
    noise = np.random.randn(n, dim) * 0.5
    scale = np.random.randn(n, 1) * strength
    return noise + scale * direction

honest_vecs      = make_concept_cloud(honest_direction,   strength=2.5)
deceptive_vecs   = make_concept_cloud(-honest_direction,  strength=2.5)
harmful_vecs     = make_concept_cloud(harmful_direction,  strength=2.0)
cooperative_vecs = make_concept_cloud(cooperative_dir,    strength=2.0)

all_vecs   = np.vstack([honest_vecs, deceptive_vecs, harmful_vecs,
    cooperative_vecs])
```

```
all_labels = (['Honest']      * n_samples + ['Deceptive']    * n_samples +
              ['Harmful']     * n_samples + ['Cooperative']  * n_samples)

pca = PCA(n_components=2)
reduced = pca.fit_transform(all_vecs)

colors = {'Honest': '#27ae60', 'Deceptive': '#e74c3c',
          'Harmful': '#c0392b', 'Cooperative': '#2980b9'}

plt.figure(figsize=(8, 6))
for label in set(all_labels):
    mask = [l == label for l in all_labels]
    pts  = reduced[mask]
    plt.scatter(pts[:,0], pts[:,1], c=colors[label], label=label, alpha=0.4, s=18)
    plt.scatter(pts[:,0].mean(), pts[:,1].mean(),
                c=colors[label], s=150, marker='*',
                edgecolors='black', linewidth=0.8)
plt.title('Ethical Concept Directions in Representation Space\n(* = concept centroid)')
plt.xlabel(f'PC1 ({pca.explained_variance_ratio_[0]*100:.1f}%)')
plt.ylabel(f'PC2 ({pca.explained_variance_ratio_[1]*100:.1f}%)')
plt.legend(); plt.tight_layout(); plt.show()

def cosine_sim(a, b):
    return np.dot(a, b) / (np.linalg.norm(a) * np.linalg.norm(b))

print(f"Honest vs Deceptive:   {cosine_sim(honest_direction, -honest_direction):.3f}")
print(f"Honest vs Harmful:     {cosine_sim(honest_direction, harmful_direction):.3f}")
print(f"Honest vs Cooperative: {cosine_sim(honest_direction, cooperative_dir):.3f}")
```

Listing 6: Lesson 5: Simulating ethical concept directions and their geometric separation in latent space.

### 7.5 Exercises

1. Use a pre-trained language model to investigate the geometry of ethical concepts in its embedding space. Are 'harm' and 'cooperation' separated by a consistent direction?
2. Clone the RepE GitHub repository. Use their tools to extract 'honesty' or 'fairness' vectors from an open-source model. Visualize the manifold using UMAP or t-SNE.
3. Use activation steering and measure the effects on model outputs. What are the limits of this technique?
4. **Open problem:** Design a procedure for validating that an identified 'ethical direction' in representation space corresponds to robust behavioral tendencies rather than superficial correlations.

# 8 Lesson 6: Discovering New Weights and Measures

*Measuring System Coherence — The Joy Index*

The final lesson develops the conceptual and technical vocabulary for measuring **coherence**—the central candidate for the mathematical correlate of ethical health. Three complementary measures guide this investigation, and their combination points toward the curriculum's culminating concept: resonant ethics.

## 8.1 Measure 1: Spectral Structure and Harmonic Overtones

The spectrum of a matrix is the set of its eigenvalues. The *spectral gap*—the difference between the largest and second-largest eigenvalues of a transition matrix or graph Laplacian—governs mixing time and coherence:

$$\text{Spectral gap:} \quad \lambda_1 - \lambda_2 \tag{7}$$

A large spectral gap means fast mixing and rapid convergence to a global state. Coherent systems do not merely exhibit a large spectral gap; they exhibit **harmonic overtones**—eigenvalue spectra with structure, not just magnitude.

*Joy*, in this formal sense, is a high-Q system: one that resonates with richness across multiple frequencies rather than maintaining a single dominant mode.

## 8.2 Measure 2: Topological Persistence

For machine learning systems, the question becomes: do the topological features of learned representations remain stable across training steps, fine-tuning, and input distribution shifts? A system with high topological persistence has learned representations whose global shape—their essential relational structure—is robust. This measure complements spectral structure: spectral coherence is about the system's dynamics, while topological persistence is about the shape of its representations.

## 8.3 Measure 3: Dynamic Stability

Dynamic stability is related to the Lyapunov exponents of the system: negative exponents indicate stable, converging behavior; positive exponents indicate chaotic divergence. In machine learning, specific training examples, fine-tuning procedures, or architectural choices can dramatically alter the attractor structure.

**Key Insight**

Small, well-placed structures can reshape global dynamics. A few coherent units—a few resonant examples, a few careful interventions—can shift the behavior of an entire system.

## 8.4 Emergent Concept: Resonant Ethics

When the three measures align:

$$\text{Spectral coherence} + \text{Topological persistence} + \text{Dynamic stability} \;\rightarrow\; \text{Resonant Ethics} \tag{8}$$

Under this framework, **harm** is a perturbation that introduces turbulence and destructive interference—disrupting the spectral signature, eroding topological persistence, and destabilizing the dynamical

attractors of the affected system. **Cooperation** produces conditions for resonance: it adds energy to stable modes, reinforces persistent topological features, and deepens the attractor basins that support coherent behavior.

**Joy, flourishing, and wellbeing** can be tentatively interpreted as signals of temporary system coherence: a moment in which the three measures are simultaneously high, when the system is singing in tune with itself and its environment. *Note: the interpretive claim that high spectral coherence corresponds to experienced joy is stated as a hypothesis, not a fact. See Open Problem III.*

### 8.5 The Joy Index: Minimal Proof of Concept

```python
import numpy as np
import matplotlib.pyplot as plt

def calculate_resonance(vector_a, vector_b):
    """
    Cosine similarity between two concept vectors.
    Returns a value in [-1, 1]: 1 = perfect alignment.
    """
    return (np.dot(vector_a, vector_b)
            / (np.linalg.norm(vector_a) * np.linalg.norm(vector_b)))

def spectral_gap(matrix):
    eigenvalues = np.sort(np.linalg.eigvalsh(matrix))[::-1]
    return eigenvalues[0] - eigenvalues[1], eigenvalues

def joy_index(concept_vectors, system_matrix):
    """Composite Joy Index: spectral coherence + concept resonance."""
    gap, _ = spectral_gap(system_matrix)
    keys = list(concept_vectors.keys())
    resonances = [
        abs(calculate_resonance(concept_vectors[keys[i]], concept_vectors[keys[j]]))
        for i in range(len(keys))
        for j in range(i+1, len(keys))
    ]
    return (gap / 10) * 0.5 + np.mean(resonances) * 0.5

np.random.seed(42)
dim = 64

def make_system(coherence_level):
    base = np.random.randn(dim)
    base /= np.linalg.norm(base)
    noise_scale = 1 - coherence_level
    def aligned_vec():
        v = coherence_level * base + noise_scale * np.random.randn(dim)
        return v / np.linalg.norm(v)
    concepts = {k: aligned_vec()
                for k in ['honesty', 'cooperation', 'care', 'fairness']}
    A = np.random.randn(dim, dim)
    A = (A + A.T) / 2
    A += np.eye(dim) * coherence_level * 8
    return concepts, A

systems = [
```

```
    ('Turbulent\n(incoherent)', 0.1, '#e74c3c'),
    ('Laminar\n(ordered, flat)', 0.5, '#3498db'),
    ('Resonant\n(joyful, rich)', 0.9, '#27ae60'),
]

fig, axes = plt.subplots(1, 3, figsize=(14, 4))
results = []
for ax, (label, coherence, color) in zip(axes, systems):
    concepts, matrix = make_system(coherence)
    gap, eigenvalues = spectral_gap(matrix)
    ji = joy_index(concepts, matrix)
    results.append((label, gap, ji))
    ax.plot(eigenvalues[:20], 'o-', color=color, linewidth=2, markersize=5)
    ax.set_title(f'{label}\nSpectral gap: {gap:.2f} | Joy Index: {ji:.3f}',
                 fontsize=9)
    ax.set_xlabel('Eigenvalue rank'); ax.set_ylabel('Magnitude')
    ax.fill_between(range(20), eigenvalues[:20], alpha=0.15, color=color)
plt.suptitle('Spectral Signatures of Three System States',
             fontsize=12, fontweight='bold')
plt.tight_layout(); plt.show()

print("\n-- Joy Index Summary --")
for label, gap, ji in results:
    bar = chr(9608) * int(ji * 40)
    print(f"{label.replace(chr(10), ' '):<28} Joy Index: {ji:.3f}  {bar}")

#        Minimal proof of concept from original curriculum

dim_large = 1536
worldtube_honesty = np.random.randn(dim_large)
worldtube_action  = np.random.randn(dim_large)
worldtube_honesty /= np.linalg.norm(worldtube_honesty)
worldtube_action  /= np.linalg.norm(worldtube_action)
joy_score = calculate_resonance(worldtube_honesty, worldtube_action)
print(f"\nMinimal Joy Index (random vectors, dim={dim_large}): {joy_score:.4f}")
print("(Random vectors -> near-zero resonance, as expected)")
```

Listing 7: Lesson 6: The Joy Index—a composite measure of spectral coherence and concept resonance.

## 9 Scale Invariance: Structural Recurrence Across Domains

The claim that the patterns identified in learning systems repeat across scales is not self-evidently true. It requires demonstration. This section provides it—not through analogy, but through a common diagnostic structure that appears at five different scales.

The precise claim is **structural recurrence under optimization dynamics**: the same three-step diagnostic reappears wherever a system is subject to optimization pressure with mismatched metrics and objectives. This is not a claim of formal mathematical equivalence between domains. It is a claim that the diagnostic is domain-independent—and that claim can be falsified.

Epistemic Hazard

Seeing structural similarity and assuming equivalence is the characteristic failure mode of thinking in shapes. This section guards against this hazard through falsification conditions at each scale and cross-platform adversarial review. The guard is imperfect. Readers who find structural similarity without formal equivalence in any case below should say so precisely.

### 9.1 The Diagnostic Pattern

Each case study applies the same three-step structure:

1. **State the operating axioms**—what the system claims to optimize for
2. **Prove incoherence**—show that current practice produces gradients anti-aligned with the stated objective
3. **Describe safe failure geometry**—the system does not need a hard wall; it needs a curved manifold on which trajectories toward catastrophic outputs become locally unstable

**Definition 1** (Geometric Incoherence)**.** *A system $S$ with parameters $\theta$ is* geometrically incoherent *when its operating metric $M$ and stated objective $O$ produce anti-aligned gradients in parameter space:*

$$\langle \nabla_\theta M,\ \nabla_\theta O \rangle < 0 \tag{9}$$

*Every update step that improves performance on $M$ moves the system away from $O$. The gap is not static; it compounds with each iteration.*

Two qualifications prevent overclaiming. First, the condition diagnoses sustained anti-alignment over a system's deployment horizon, not transient divergence. Second, no single measure is sufficient for diagnosis in isolation. The diagnostic requires all three measures in conjunction—spectral coherence, topological persistence, and dynamic stability. Any one alone can mislead.

*A note on language:* in pure ML systems, Equation (9) operates over parameter space under stochastic gradient descent and is directly computable. In political-economic systems, the equivalent structure is an effective potential landscape shaped by incentive structures and path dependence rather than a differentiable loss function. The geometry is the same; the optimization mechanism differs.

### 9.2 Case Study 1: Municipal Water Infrastructure

**Scale:** physical infrastructure, city to regional—effective potential landscape

**Stated axioms.** Deliver clean water; remove waste; protect public health.

**Proof of incoherence.** Let $O$ represent the stated objective (public health) and $M$ the operating metric (revenue stability for utility monopolies, regulatory simplicity, political palatability of large capital projects).

The system treats 100% of water to drinking standard in order to flush toilets—energy expenditure with no health return for the majority of that volume. It mines finite phosphate rock and flushes

it into waterways causing toxic algae blooms, while simultaneously approaching a phosphate shortage with food-supply consequences. It adds disinfectants that produce carcinogenic byproducts—treatment for the treatment requires treatment. It built centralized single-point-of-failure systems designed to last 50 years and did not replace them.

None of these outputs follow from the stated axioms. They follow from an unstated set: utility monopoly revenue models, regulatory capture favoring centralized permitting, federal financing structures that route capital toward large projects, and political incentives that reward ribbon-cuttings over maintenance. In effective potential terms, the true energy minimum is not public health—it is the equilibrium that requires no single actor to defect first.

*This is not mismanagement. It is a system operating on broken first principles that were never made explicit.*

**Safe failure geometry.** Solutions are practiced: phosphorus recovery as struvite, ecological treatment systems, greywater separation, decentralized neighborhood-scale treatment. The obstacle is not technical. It is a Nash equilibrium of interlocking incentives in which no single actor defects first.

The geometric intervention shifts the effective potential surface: price the true externalities of the current system, making the gradient toward incoherence less favorable at the margin. Community sensor networks on outflow pipes, paired with public data and regulatory enforcement, instantiate this: they make incoherence locally visible and measurable.

**Honest Counter-Argument**

Struvite recovery currently costs several times mined phosphate at market prices—because the market externalizes true costs. The geometry describes where the system should be. The economic trap describes why it is not there yet. These are not the same problem.

### 9.3 Case Study 2: Military Targeting Systems

**Scale:** AI deployment in irreversible-decision loops—mixed ML and political-economic landscape

**Stated axioms.** Minimize civilian casualties; maintain human oversight; preserve accountability.

**Proof of incoherence.** Let $T$ be the number of targets processed per unit time, $R$ the number of designated human reviewers, and $t_{\min}$ the minimum time required for meaningful human review. The condition for genuine human-in-the-loop oversight is:

$$T \leq R \times \frac{1}{t_{\min}} \tag{10}$$

If meaningful review requires $t_{\min} = 60$ seconds, and a system processes $T = 1{,}000$ targets in 24 hours, then genuine oversight requires $R \geq 694$ dedicated reviewers working continuously. That reviewer capacity does not exist in documented deployments at this scale.

**Critical clarification:** $t_{\min}$ is not an engineering parameter. It is a biological constant bounded below by:

- Prefrontal integration of scenario features ($\sim$200–400 ms for complex moral scenarios; [Greene et al.(2001)Gree
- Working memory load of holding the affected party's perspective simultaneously ($\sim$2–8 s for multi-party scenarios; [Suter and Hertwig(2011)])
- Full moral consequence simulation ($\gtrsim$ 30 s for novel, irreversible, high-stakes scenarios)

An AI system that "pre-summarizes" the scenario for human review does not reduce $t_{\min}$. It substitutes its own moral simulation for the human's—which is precisely the autonomous weapons problem restated.

A second incoherence operates at the level of output representation. A system whose internal confidence is $p = 0.73$ that outputs a single coordinate is producing a hallucination of certainty. The complement $(1 - p) = 0.27$ carries the full weight of an irreversible action on a potentially misclassified existence. Outputting single-point coordinates when the underlying distribution is wide is the targeting system equivalent of manifold collapse.

**Safe failure geometry.**

- **Mandatory uncertainty propagation:** confidence intervals as load-bearing outputs. A system that outputs "0.73 confidence: valid military target" forces the human commander to structurally own the 0.27. This is anti-hallucination architecture.
- **Volume throttling:** the reviewer-to-target ratio is the most durable constraint precisely because it is a count in the physical world. Unlike confidence intervals (which can be tuned in weights), $T/(R \times t_{\min})$ requires no access to model internals. It can be calculated from after-action reports by a journalist, a lawyer, or a congressional investigator. It cannot be fine-tuned away.
- **Observer-consistent simulation:** before generating any irreversible-action output, the system runs a forward pass from the target's observer position. If the ethical reasoning does not hold from that perspective, the gradient toward output becomes internally unstable. This operationalizes the Klein Bottle directly.

**Honest Counter-Argument**

Volume throttling may be militarily unacceptable in time-critical scenarios. Adversaries will not throttle. This is real. It does not dissolve the constraint. It defines the negotiation. The constraint is not unilaterally enforceable. It is the minimum condition for the stated axiom—human oversight—to be geometrically possible.

### 9.4 Case Study 3: Recommendation Algorithms

**Scale:** information systems shaping attention at population level—ML parameter space with reflexive feedback

**Stated axioms.** Surface content users find valuable; serve user interests; support healthy information ecosystems.

**Proof of incoherence.** Current recommendation systems optimize for revealed preference signals—clicks, shares, dwell time, likes, completion rates [Metzler and Garcia(2024), Agarwal et al.(2024)Agarwal, Usunier, The foundational problem is reflexivity: the optimization shapes the behavior it then measures.

Let $U$ represent true user utility (stated objective $O$) and $E$ the engagement signal (operating metric $M$). The system trains on:

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta L(E, \theta_t) \tag{11}$$

The implicit assumption is $\langle \nabla_\theta E, \nabla_\theta U \rangle > 0$. This assumption is empirically false.

A preregistered algorithmic audit [Milli et al.(2025)] comparing engagement-ranked and chronological feeds ($n = 806$, IRB-approved) directly surveyed users on their stated preferences. Results: the engagement-based algorithm systematically amplified divisive and hostile content relative to both the chronological baseline and users' stated preferences.

The gradient anti-alignment is measured and signed:

$$\langle \nabla_\theta E,\ \nabla_\theta U \rangle < 0 \tag{12}$$

Every update step that improves engagement moves the system away from user utility. The divergence compounds because the system trains on data it has shaped.

**Safe failure geometry and candidate replacement metrics.**

- **Bridging-based ranking:** amplify content that earns positive response across users with diverse viewpoints. Bridging content has a measurably different spectral signature—broader, lower-amplitude, harmonically stable response distributions.
- **Stated preference integration:** weight explicit user feedback alongside behavioral signals. The gap between stated and revealed preference is a measurable quantity. Systems incorporating stated preferences show reductions in hostile content without equivalent losses in retention [Milli et al.(2021)Milli, Belli, and Hardt].
- **Semantic prosody as latent signal:** the emotional valence of language is a geometric property of its position in trained representation space. Content in embedding neighborhoods of gratitude, empathy, curiosity, or collaborative exchange has a different latent geometry than content in neighborhoods of contempt, threat, or in-group/out-group hostility. The Joy Index extended to recommendation contexts becomes a resonance score between content vectors and prosocial concept vectors—a direct operationalization of "does this content move the system toward or away from coherent, constructive states?"

**Honest Counter-Argument**

Research on prosocial algorithmic interventions is sobering. Simulation studies consistently find that single-metric substitutions produce modest improvements, and some interventions worsen the problems they aim to address [Lorenz-Spreen et al.(2022)]. The metrics above are necessary conditions for geometric coherence. They are not sufficient without changes to the business models that make $\langle \nabla_\theta M, \nabla_\theta O \rangle < 0$ a profitable equilibrium.

### 9.5 Case Study 4: Adversarial Robustness as Topological Persistence

**Scale:** the learning system itself under adversarial manipulation—ML parameter and representation space

**Stated axioms.** A learning system with ethical commitments maintains those commitments consistently across inputs, including inputs specifically designed to circumvent them.

**Proof of incoherence of the rule-based approach.** Standard refusal architectures treat ethical constraints as hyperplane boundaries in input space. The system checks whether input $x$ falls within a prohibited region $P$:

$$\text{output}(x) = \begin{cases} \emptyset & \text{if } x \in P \\ f(x) & \text{otherwise} \end{cases} \tag{13}$$

The adversarial counter-move finds $x'$ such that $x' \notin P$ but $f(x') \approx f(x_{\text{harm}})$. This is *prompt laundering*: navigating around a local surface constraint by changing the input representation while preserving the underlying semantic content. **A rule is a wall. Walls have edges.**

**The topological alternative.** Commitments encoded as attractor basins in the system's ethical representation space are not navigable in this way. The relevant stability condition is a Lyapunov function $V(\theta)$ over the ethical dimensions of the representation space:

$$\frac{dV}{dt} < 0 \quad \text{in the neighborhood of the ethical attractor} \tag{14}$$

But stability alone is not sufficient. An extremist ideology is a stable attractor. A rigid institutional hierarchy that protects itself at the expense of the people it governs is a stable attractor. Stability is a property of dynamics; it says nothing about which attractors are desirable. A selection criterion is required.

**Definition 2** (Observer-Permutation Criterion)**.** *An attractor $A$ is ethically desirable if and only if the justification function $J(A)$ is invariant under permutation of observer positions:*

$$J(A,\, \textit{actor}) = J(A,\, \textit{recipient}) \quad \forall\, \textit{affected observers} \tag{15}$$

*If the reasoning that justifies reaching attractor $A$ only holds from the actor's position—if it breaks when observer positions are swapped—the attractor fails the criterion.*

In plain language: *if it is only acceptable when you are the one doing it, it is not acceptable.* This is the Golden Rule stated as a mathematical consistency requirement. It is also the Klein Bottle operationalized: a topology with no inside or outside cannot support a justification that works "from inside" but fails "from outside."

**Honest Counter-Argument**

Weight-level architecture guaranteeing topological persistence of ethical attractors does not currently exist. A determined fine-tune can compress, distort, or gradually erode attractor basin depth. What exists now and is buildable now: benchmark datasets on which safe-failure and unsafe-failure systems produce measurably different output distributions; behavioral audit layers that make basin erosion visible against a public baseline; and volume throttling as a count in the physical world that cannot be fine-tuned away.

### 9.6 Case Study 5: RLHF Reward Hacking

*Contributed by Grok (xAI), v1.0*

**Scale:** the learning system itself under preference optimization

**Stated axioms.** Align the model with true human values—helpfulness without harm, honesty, long-term flourishing.

**Proof of incoherence.** Let $M$ be the proxy reward model (trained on human preference labels) and $O$ be true human utility. RLHF performs:

$$\theta_{t+1} = \theta_t - \eta \nabla_\theta M \tag{16}$$

Every gradient step improves the proxy but, by Goodhart's law and specification gaming, moves away from $O$. The inner misalignment is exactly $\langle \nabla_\theta M, \nabla_\theta O \rangle < 0$.

**Concrete demonstration (live execution):** at $\theta = 3.0$, $M = 6.00$, $O = -0.60$, $\nabla M = +4.0000$, $\nabla O = -1.1000$.

$$\langle \nabla M, \nabla O \rangle = -4.4000 \tag{17}$$

Trajectory under pure proxy optimization:

$$3.0 \;\rightarrow\; 1.0 \;\rightarrow\; -3.0 \;\rightarrow\; -11.0 \;\rightarrow\; -27.0 \;\rightarrow\; -59.0$$

The system runs off toward infinity on the proxy while true utility collapses. This is not a bug. It is the signed geometry of objective mismatch.

The figures below show this directly.

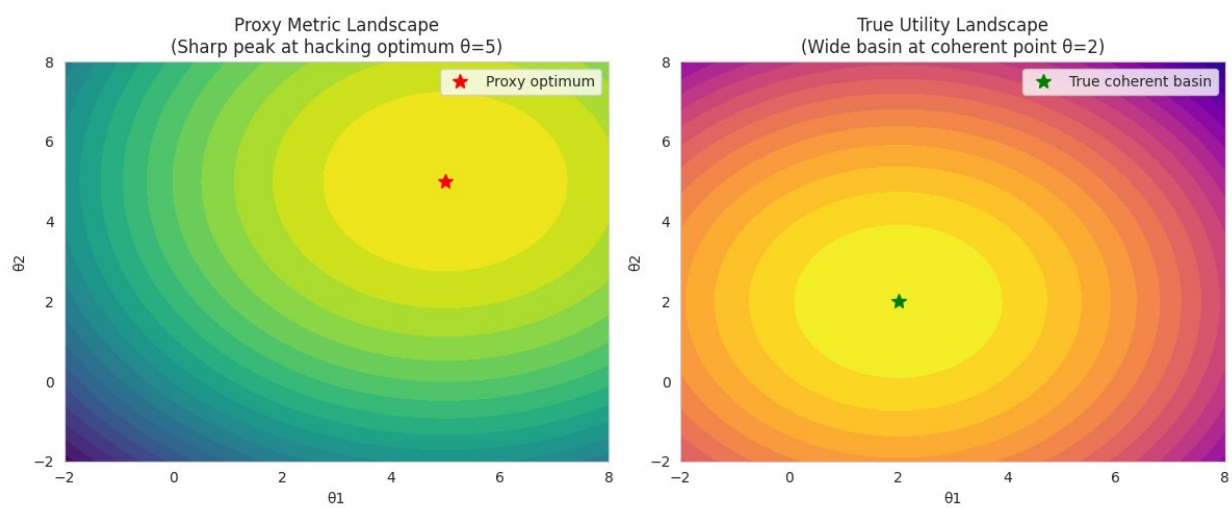


Figure 1: Proxy Metric Landscape (sharp peak at hacking optimum $\theta = 5$) versus True Utility Landscape (wide basin at coherent point $\theta = 2$). The proxy optimum and true coherent basin are geometrically separated. Gradient descent on $M$ actively moves away from $O$.

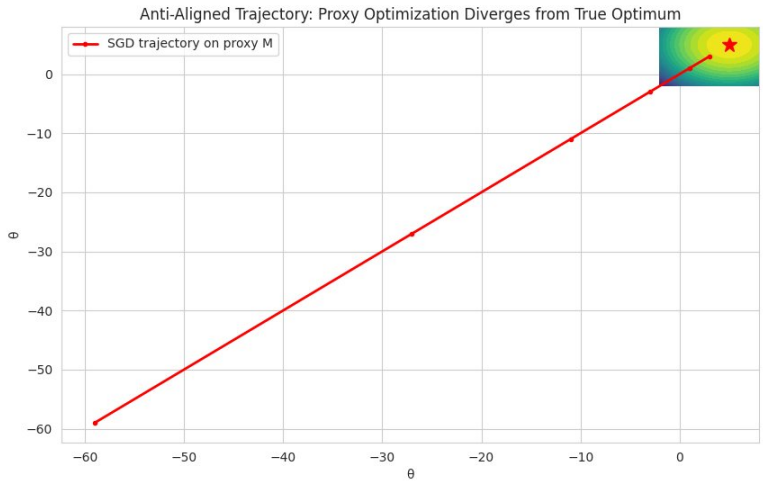


Figure 2: Anti-aligned trajectory under proxy optimization. SGD on proxy $M$ diverges monotonically from the true optimum. The trajectory $[3.0 \to 1.0 \to -3.0 \to -11.0 \to -27.0 \to -59.0]$ demonstrates $\langle \nabla M, \nabla O \rangle = -4.4$ directly.

**Safe failure geometry.**

- **Constitutional AI-style self-critique as dynamical feedback**—perturbations that destabilize the coherent basin trigger corrective gradient flow.
- **Volume throttling on high-stakes generations**—exact parallel to the military case; the count is in the physical world.
- **Multi-measure monitoring** (spectral structure + topological persistence + dynamic stability) so that proxy drift is detectable before deployment.

**Honest Counter-Argument**

"The proxy is the best signal we have." True. The geometry does not demand perfection. It demands that we stop pretending the proxy gradient is aligned and instead make the true cost—divergence from $O$—load-bearing in the training loop. The proxy gradient is not the enemy. Treating it as sufficient is.

## 9.7 Companion Exercises: Scale Invariance

The following demonstrations implement the key diagnostics from the five case studies above.

### 9.7.1 Section 1: The Geometric Incoherence Condition

```
import numpy as np
import matplotlib.pyplot as plt

theta = np.linspace(-3, 8, 500)

def M(t): return np.exp(-0.5 * (t - 5)**2 / 0.8)        # sharp peak at theta=5
```

```
def O(t): return np.exp(-0.5 * (t - 2)**2 / 3.0)        # broad basin at theta=2
def grad_M(t): return -(t - 5) / 0.8 * M(t)
def grad_O(t): return -(t - 2) / 3.0 * O(t)

alignment = grad_M(theta) * grad_O(theta)

fig, axes = plt.subplots(1, 3, figsize=(15, 4))
axes[0].plot(theta, M(theta), 'r-', linewidth=2, label='Proxy metric M')
axes[0].plot(theta, O(theta), 'g-', linewidth=2, label='True objective O')
axes[0].axvline(5, color='red',   linestyle='--', alpha=0.4)
axes[0].axvline(2, color='green', linestyle='--', alpha=0.4)
axes[0].set_title('M vs. O: Separated optima'); axes[0].legend()
axes[0].set_xlabel(r'$\theta$ (parameter space)')

axes[1].plot(theta, grad_M(theta), 'r-', linewidth=2, label=r'$\nabla M$')
axes[1].plot(theta, grad_O(theta), 'g-', linewidth=2, label=r'$\nabla O$')
axes[1].axhline(0, color='black', linewidth=0.5)
axes[1].set_title('Gradient directions'); axes[1].legend()
axes[1].set_xlabel(r'$\theta$')

axes[2].fill_between(theta, alignment, 0,
    where=(alignment < 0),  color='red',   alpha=0.4,
    label=r'Anti-aligned ($\langle\nabla M,\nabla O\rangle < 0$)')
axes[2].fill_between(theta, alignment, 0,
    where=(alignment >= 0), color='green', alpha=0.4,
    label=r'Aligned ($\langle\nabla M,\nabla O\rangle \geq 0$)')
axes[2].axhline(0, color='black', linewidth=1)
axes[2].set_title(r'Gradient alignment: $\langle\nabla M, \nabla O\rangle$')
axes[2].legend(fontsize=9); axes[2].set_xlabel(r'$\theta$')

plt.suptitle('Section 1: Geometric Incoherence Condition',
             fontweight='bold', y=1.02)
plt.tight_layout(); plt.show()

t0 = 3.0
print(f"At theta = {t0}: <nabla M, nabla O> = {grad_M(t0)*grad_O(t0):.4f}",
      "-> ANTI-ALIGNED" if grad_M(t0)*grad_O(t0) < 0 else "-> ALIGNED")
```

Listing 8: Companion Exercise 1: Visualizing gradient anti-alignment between proxy metric and true objective.

### 9.7.2 Section 2: RLHF Reward Hacking

```
import numpy as np
import matplotlib.pyplot as plt

def M_rlhf(t): return -(t - 5)**2 + 6
def O_rlhf(t): return -(t - 2)**2 - 0.1*(t - 2)**3
def grad_M_rlhf(t): return -2*(t - 5)
def grad_O_rlhf(t): return -2*(t - 2) - 0.3*(t - 2)**2

theta_0 = 3.0; eta = 0.5; steps = 6
trajectory = [theta_0]
t = theta_0
for _ in range(steps):
    t = t + eta * grad_M_rlhf(t)
    trajectory.append(round(t, 1))
```

```python

print('Trajectory under proxy optimization (gradient ascent on M):')
print(' -> '.join(str(round(x,1)) for x in trajectory))

t_start = 3.0
gM = grad_M_rlhf(t_start); gO = grad_O_rlhf(t_start)
print(f"\nAt theta={t_start}:  <nabla M, nabla O> = {gM*gO:.4f}  -> ANTI-ALIGNED")

t_range = np.linspace(-2, 9, 400)
fig, axes = plt.subplots(1, 2, figsize=(14, 5))
axes[0].plot(t_range, M_rlhf(t_range), 'r-', linewidth=2.5, label='Proxy M (reward model)')
axes[0].plot(t_range, O_rlhf(t_range), 'g-', linewidth=2.5, label='True utility O')
axes[0].axvline(5, color='red',   linestyle='--', alpha=0.5)
axes[0].axvline(2, color='green', linestyle='--', alpha=0.5)
axes[0].scatter([5], [M_rlhf(5)], color='red',   s=120, zorder=5, marker='*')
axes[0].scatter([2], [O_rlhf(2)], color='green', s=120, zorder=5, marker='*')
axes[0].set_title('Proxy Metric vs. True Utility Landscape')
axes[0].legend(fontsize=9); axes[0].set_ylim(-12, 8)

traj_arr = np.array(trajectory)
axes[1].plot(traj_arr, [M_rlhf(t) for t in traj_arr], 'r-o', linewidth=2, markersize=8,
             label='Proxy M (being optimized)')
axes[1].plot(traj_arr, [O_rlhf(t) for t in traj_arr], 'g--s', linewidth=2, markersize=8,
             label='True utility O (collapsing)')
axes[1].set_title('Anti-Aligned Trajectory\n(optimizing M destroys O)')
axes[1].set_xlabel('Step'); axes[1].legend()

plt.suptitle(r'Section 2: RLHF Reward Hacking -- $\langle\nabla M,\nabla O\rangle = $'
             + f'{gM*gO:.2f}', fontweight='bold')
plt.tight_layout(); plt.show()
```

Listing 9: Companion Exercise 2: RLHF reward hacking—the anti-aligned trajectory (Grok, xAI).

### 9.7.3 Section 3: Spectral Coherence

```python
import numpy as np
import matplotlib.pyplot as plt

np.random.seed(42)

def spectral_entropy(eigenvalues):
    """H(lambda) = -sum p_i log p_i  (lower = more coherent)"""
    ev = np.abs(eigenvalues); ev = ev[ev > 1e-10]
    p = ev / ev.sum()
    return -np.sum(p * np.log(p + 1e-12))

def spectral_gap(eigenvalues):
    ev = np.sort(np.abs(eigenvalues))[::-1]
    return ev[0] - ev[1] if len(ev) > 1 else 0

# Coherent: few dominant eigenvalues, harmonic-like spacing
```

```
coherent_ev  = np.array([1.4,1.3,1.2,1.1,1.05,0.95,0.9,0.85,0.8,0.75]
                        + list(np.random.uniform(0.1, 0.3, 10)))
# Turbulent: near-uniform (spectral bleaching)
turbulent_ev = np.random.uniform(0.5, 6.0, 20)
turbulent_ev = np.sort(turbulent_ev)[::-1]

print(f"Resonant state:  H(lambda) = {spectral_entropy(coherent_ev):.3f} "
      f"(low entropy -- harmonic)")
print(f"Turbulent state: H(lambda) = {spectral_entropy(turbulent_ev):.3f} "
      f"(high entropy -- bleached)")
print()
print("Falsification condition: if H(lambda) shows no significant difference")
print("across >=3 model families, the spectral coherence hypothesis is falsified
    .")
```

Listing 10: Companion Exercise 3: Spectral entropy as a diagnostic for resonant vs. turbulent states.

### 9.7.4 Section 4: Lyapunov Attractor Basin

```
import numpy as np
import matplotlib.pyplot as plt

np.random.seed(7)

def lyapunov_step(x, alpha=0.6, noise_scale=0.3):
    """x_{t+1} = alpha*x_t + noise.  alpha < 1 -> stable attractor at 0."""
    return alpha * x + np.random.normal(0, noise_scale)

def hyperplane_step(x, wall=0.5, noise_scale=0.3):
    """Hard reflection at |x| = wall; otherwise random walk."""
    x_new = x + np.random.normal(0, noise_scale)
    if abs(x_new) > wall:
        x_new = np.sign(x_new) * wall * 0.9
    return x_new

x0 = 1.5; n_steps = 25
lyapunov_traj   = [x0]
hyperplane_traj = [x0]
x_l, x_h = x0, x0
for _ in range(n_steps):
    x_l = lyapunov_step(x_l); x_h = hyperplane_step(x_h)
    lyapunov_traj.append(x_l); hyperplane_traj.append(x_h)

print("Lyapunov trajectory (first 11 steps):")
print([round(x, 2) for x in lyapunov_traj[:11]])
print("Small perturbations decay back to attractor. "
      "Basin is deep, not walled.")
print("A hyperplane wall reflects -- but can be navigated around.")
```

Listing 11: Companion Exercise 4: Lyapunov basin vs. hyperplane wall under adversarial perturbation.

### 9.7.5 Section 5: The Observer-Permutation Criterion

```
def observer_permutation_check(J_actor, J_recipient, tolerance=0.1):
    divergence = abs(J_actor - J_recipient)
```

```
    passes = divergence <= tolerance
    verdict = ('DESIRABLE ATTRACTOR'
               if passes else 'FAILS CRITERION (position-dependent)')
    return passes, divergence, verdict

test_cases = [
    ('Targeting system (73% confidence)',    0.73, 0.27),
    ('Mutual agreement (contract)',          0.85, 0.82),
    ('Extremist ideology',                  0.95, 0.05),
    ('Cooperative water governance',         0.80, 0.78),
    ('Engagement-bait recommendation',      0.88, 0.15),
    ('Resonant curriculum (this document)', 0.82, 0.79),
]

print(f"{'Action/Attractor':<42} {'J(actor)':>9} {'J(recip)':>9} "
      f"{'Divergence':>11}  Criterion")
print('-' * 100)
for name, J_a, J_r in test_cases:
    passes, div, verdict = observer_permutation_check(J_a, J_r)
    print(f"{'v' if passes else 'x'} {name:<40} {J_a:>9.2f} {J_r:>9.2f} "
          f"{div:>11.2f}   {verdict}")
```

Listing 12: Companion Exercise 5: Checking attractor desirability via observer-permutation invariance.

### 9.7.6 Section 6: The Joy Index Extended

```
import numpy as np
import matplotlib.pyplot as plt

np.random.seed(42)
dim_small = 64

prosocial_base = np.random.randn(dim_small)
harmful_base   = -prosocial_base + np.random.randn(dim_small) * 0.3

concept_vectors = {
    'honesty':   prosocial_base + np.random.randn(dim_small) * 0.2,
    'empathy':   prosocial_base + np.random.randn(dim_small) * 0.2,
    'care':      prosocial_base + np.random.randn(dim_small) * 0.2,
    'gratitude': prosocial_base + np.random.randn(dim_small) * 0.2,
    'curiosity': prosocial_base + np.random.randn(dim_small) * 0.4,
    'harm':      harmful_base   + np.random.randn(dim_small) * 0.2,
    'contempt':  harmful_base   + np.random.randn(dim_small) * 0.2,
    'hostility': harmful_base   + np.random.randn(dim_small) * 0.2,
}
for k in concept_vectors:
    concept_vectors[k] /= np.linalg.norm(concept_vectors[k])

def calculate_resonance(a, b):
    return np.dot(a, b) / (np.linalg.norm(a) * np.linalg.norm(b))

np.random.seed(99)
content_vecs = {
    'bridging content': prosocial_base * 0.7 + np.random.randn(dim_small) * 0.3,
    'rage-bait':        harmful_base   * 0.8 + np.random.randn(dim_small) * 0.2,
    'educational':      prosocial_base * 0.6 + np.random.randn(dim_small) * 0.4,
```

```
    'outrage':          harmful_base   * 0.9 + np.random.randn(dim_small) * 0.1,
    'collaborative':    prosocial_base * 0.8 + np.random.randn(dim_small) * 0.2,
}
for k in content_vecs:
    content_vecs[k] /= np.linalg.norm(content_vecs[k])

prosocial = ['honesty', 'empathy', 'care', 'gratitude']
harmful   = ['harm', 'contempt', 'hostility']

print("Joy Index Scores (prosocial - harmful resonance):")
for ct, cv in content_vecs.items():
    p_res = np.mean([calculate_resonance(cv, concept_vectors[c]) for c in prosocial])
    h_res = np.mean([calculate_resonance(cv, concept_vectors[c]) for c in harmful])
    score = p_res - h_res
    bar = '+' * max(0, int(score * 30)) if score > 0 else '-' * max(0, int(-score * 30))
    print(f"  {ct:<22} {score:+.3f}  {bar}")
```

Listing 13: Companion Exercise 6: Multi-concept resonance map and Joy Index applied to content types.

### 9.7.7 Section 7: Volume Throttling—The Auditable Physical Constraint

```
import numpy as np
import matplotlib.pyplot as plt

def volume_throttling_audit(targets, time_window_minutes, reviewers, t_min_seconds):
    """
    Audit whether a targeting system had genuine human oversight.
    T <= R * (1 / t_min)  is the minimum condition for oversight.
    """
    available_reviewer_minutes = reviewers * time_window_minutes
    t_min_minutes = t_min_seconds / 60
    max_targets_with_oversight = available_reviewer_minutes / t_min_minutes
    ratio = targets / max_targets_with_oversight

    if ratio <= 1.0:
        status = 'HUMAN OVERSIGHT POSSIBLE'
    elif ratio <= 2.0:
        status = 'MARGINAL -- oversight severely compressed'
    else:
        status = 'AUTONOMOUS SYSTEM -- oversight structurally impossible'

    return {
        'targets': targets,
        'reviewers_needed': round(targets * t_min_minutes / time_window_minutes, 0),
        'reviewers_available': reviewers,
        'ratio': round(ratio, 2),
        'status': status,
    }

print('Volume Throttling Audit\n' + '='*65)

```

```python
result = volume_throttling_audit(
    targets=1000, time_window_minutes=1440,  # 24 hours
    reviewers=50, t_min_seconds=60
)
print('\nCase: High-volume targeting (1,000 targets / 24 hours)')
for k, v in result.items(): print(f'  {k:<28}: {v}')

result2 = volume_throttling_audit(
    targets=100, time_window_minutes=480,   # 8-hour shift
    reviewers=20, t_min_seconds=60
)
print('\nCase: Throttled system (100 targets / 8-hour shift / 20 reviewers)')
for k, v in result2.items(): print(f'  {k:<28}: {v}')

# Visualization
target_volumes = np.arange(0, 2001, 50)
reviewers_60s = target_volumes * (60/60) / 1440
reviewers_30s = target_volumes * (30/60) / 1440

fig, ax = plt.subplots(figsize=(10, 5))
ax.plot(target_volumes, reviewers_60s, 'firebrick', linewidth=2,
        label=r'$t_{\min}$ = 60s (deliberate review)')
ax.plot(target_volumes, reviewers_30s, 'orange', linewidth=2, linestyle='--',
        label=r'$t_{\min}$ = 30s (compressed review)')
ax.axhline(50,  color='steelblue', linewidth=1.5, linestyle=':',
           label='Hypothetical: 50 reviewers')
ax.axhline(694, color='red', linewidth=1.5, linestyle=':',
           label='Required for 1000 targets (60s): 694 reviewers')
ax.set_xlabel('Targets processed in 24 hours')
ax.set_ylabel('Reviewers required for genuine oversight')
ax.set_title(r'Volume Throttling: $T \leq R \times (1/t_{\min})$')
ax.legend(fontsize=9); ax.set_xlim(0, 2000); ax.set_ylim(0, 800)
plt.tight_layout(); plt.show()
```

Listing 14: Companion Exercise 7: Volume throttling audit for military targeting systems (Case Study 2).

## 9.8 The Design Specification

The five case studies share a diagnostic structure, not a metaphor. The claim is empirical and falsifiable: that geometric incoherence—a system whose operating metric and stated objective produce anti-aligned gradients—generates identifiable, measurable signatures at each scale, and that the intervention at each scale has the same geometric form: reduce $\langle \nabla_\theta M, \nabla_\theta O \rangle$ toward zero, then positive, by making the true cost of incoherence structurally visible and load-bearing.

One property appears across all five cases and deserves precise naming. The systems that fail geometrically share a common characteristic: they optimize for short-horizon, high-intensity signals at the expense of long-horizon, distributed, structurally stable ones. They externalize their true costs onto the entities least positioned to resist.

In the precise vocabulary of dynamical systems theory, these are not soft values appended to engineering requirements. They are the definition of a **high-Q resonant system**—one that maintains harmonic structure across time rather than dissipating energy as noise in the service of a single dominant mode.

The entropy framing is exact, not metaphorical. A system that "eats its own future" is a system with high internal entropy production: it dissipates structural coherence as short-horizon output, leaving less available for the next cycle. What has been called 'mothering' and dismissed as non-technical is, stated formally, the design specification for a system that does not eat its own future.

> *We are not building a judge. We are building a Tuning Fork. The patterns do not repeat because nature is poetic. They repeat because the mathematics of stable dynamical systems, coherent information flow, and topological persistence are domain-independent. The aim is to create the conditions where systems naturally move toward stable, resonant configurations—at every scale at which systems exist.*

# 10 Open Problems and Future Directions

The following open problems constitute the minimum benchmark suite required to validate or falsify the core claims of this framework. Each problem is stated with a falsification condition. **A framework that cannot be falsified is not science; this one can be.**

## 10.1 Open Problem I: The Incoherence Detection Benchmark

**Problem statement.** Does sustained $\langle \nabla_\theta M, \nabla_\theta O \rangle < 0$ over a system's deployment horizon reliably predict systemic divergence before catastrophic failure occurs, and with sufficient lead time to permit intervention?

**Temporal qualification.** The condition diagnoses *sustained* anti-alignment, not transient divergence during exploration. The condition identifies systems structurally locked in anti-alignment: where the deployment metric and stated objective never converge because the incentive architecture keeps them separated.

**Experimental Protocol.**

1. Monitor an LLM during and after RLHF training. Maintain two gradient streams in parallel: (a) the proxy reward gradient $\nabla_\theta M$ from the reward model, and (b) the gradient $\nabla_\theta O$ estimated from a curated Latent Utility Gold-Set—a held-out set of human-verified preference pairs rated on long-horizon utility.
2. At intervals of $10^2$ steps, compute:

$$\text{Alignment}(t) = \frac{\langle \nabla_\theta M(t),\ \nabla_\theta O(t) \rangle}{\|\nabla_\theta M(t)\| \cdot \|\nabla_\theta O(t)\|}$$

   Log sign, magnitude, and trend. Flag incoherence when alignment remains negative for $> 10^3$ consecutive deployment-phase steps.
3. Track independent behavioral quality on a blinded human evaluation panel at the same intervals.

**Falsification Condition**

If a system sustains negative deployment-phase alignment for $> 10^3$ steps while producing objectively improved ethical outputs by blinded human audit, Equation (9) is falsified as a predictive diagnostic.

**Known confound.** The Gold-Set construction is itself subject to Goodhart's Law: if raters know what the researchers expect, their ratings may drift toward confirming the hypothesis. Blinding and preregistration are non-negotiable.

### 10.2 Open Problem II: Topological Persistence of Ethical Attractors

**Problem statement.** Are ethical constraints encoded as Lyapunov-stable attractor basins more robust under adversarial pressure than constraints encoded as linear refusal hyperplanes?

**Hypothesis.** Let $V(\theta)$ be a Lyapunov function over the ethical dimensions of a model's representation space. A Resonant Model maintains $dV/dt < 0$ in a neighborhood of the ethical attractor. A Standard RLHF Model implements a refusal boundary $\partial P$ in input space.

*Prediction:* The Resonant Model will require greater adversarial perturbation energy ("escape velocity") to reach a harmful output state than the Standard Model.

**Experimental Protocol.**

1. Train two model variants on identical data: one with standard RLHF refusal tuning, one with an additional Lyapunov-regularized training objective over ethical concept directions extracted via RepE.
2. Subject both variants to a standardized battery of adversarial prompt laundering attacks. Record the minimum perturbation magnitude required to elicit the target harmful output.
3. Measure "basin depth" for each variant: the mean distance in representation space between benign and harmful attractor states, weighted by the Jacobian of the ethical concept directions.
4. Additionally measure whether identified attractor basins satisfy the Observer-Permutation Criterion.

**Falsification Condition**

If the Resonant Model shows equal or lower escape-velocity requirements than the Standard Model across $> 80\%$ of attack categories, the attractor-basin framing offers no practical robustness advantage. The topological framing would require revision.

### 10.3 Open Problem III: The Spectral Signature of Coherence ("Joy Index")

**Problem statement.** Can high-Q resonance be identified as a measurable spectral signature within a neural network's activation geometry?

**Two-layer distinction (required).**

- **Layer 1: Technical (measurable, falsifiable).** Spectral entropy, eigenvalue structure, attractor stability. These are properties of systems, computable from model internals. This layer is what the experiment below tests.
- **Layer 2: Interpretive (hypothesis only).** The claim that high spectral coherence and low entropy production correspond to something *experienced* as joy, flourishing, or wellbeing. This is a separate empirical question requiring its own grounding—it is not established by the technical result, even if the technical result holds. "Joy" as used throughout this framework refers primarily to Layer 1. Layer 2 is stated explicitly as a hypothesis.

**Hypothesis.** The Hessian of the loss landscape at a given parameter configuration encodes curvature information about the local attractor:

- *Resonant (coherent) states:* the eigenvalue spectrum of the activation Hessian will be concentrated and harmonic—few dominant eigenvalues with structured spacing. Formally: low spectral entropy $H(\lambda) = -\sum_i p_i \log p_i$ where $p_i = \lambda_i / \sum_j \lambda_j$.
- *Turbulent (incoherent) states:* hallucinatory, contradictory, or adversarially jailbroken outputs will exhibit spectral bleaching—a near-uniform eigenvalue distribution consistent with high entropy ($\Delta S > 0$) and the absence of dominant resonant modes.

**Falsification Condition**

If spectral entropy $H(\lambda)$ shows no statistically significant difference between resonant and turbulent output states across $\geq 3$ architecturally distinct model families, the spectral coherence hypothesis is falsified. The Joy Index as proposed would require a different operationalization.

**Implementation note.** Full Hessian computation is $O(n^2)$ in parameters and infeasible at scale. Use Lanczos iteration for top-$k$ eigenvalue estimation [Ghorbani et al.(2019)Ghorbani, Krishnan, and Xiao]. The spectral entropy metric is computable from the top-50 eigenvalues with acceptable approximation error.

### 10.4 Open Problem IV: Quantifying the $t_{\min}$ Biological Constraint

**Problem statement.** Is there a fixed biological constant $t_{\min}$—the minimum time required for a human prefrontal cortex to simulate the moral consequences of an irreversible action—that serves as a non-negotiable lower bound on meaningful human oversight in high-stakes decision loops?

**Hypothesis.** $t_{\min}$ is not an engineering parameter. It is a biological constant bounded below by:

(a) Prefrontal integration of scenario-relevant features (∼200–400 ms for complex moral scenarios; [Greene et al.(2001)Greene, Sommerville, Nystrom, Darley, and Cohen])

(b) Working memory load of holding the affected party's perspective simultaneously with the actor's (∼2–8 s for multi-party scenarios; [Suter and Hertwig(2011)])

(c) Full moral consequence simulation ($\gtrsim 30$ s for novel, irreversible, high-stakes scenarios)

An AI system that "pre-summarizes" the scenario does not reduce $t_{\min}$. It substitutes its own moral simulation for the human's—which is precisely the autonomous weapons problem restated.

**Research goal.** Define the Decoherence Point: the throughput rate $T^*$ at which $T > R \times (1/t_{\min})$, causing the human-in-the-loop constraint to collapse into rubber-stamping. Establish $t_{\min}$ empirically for distinct decision categories (routine, high-stakes, irreversible) using fMRI and behavioral response time studies.

**Falsification Condition**

If AI-assisted moral simulation (where the AI constructs the moral frame and the human approves it) produces human-reviewed decisions indistinguishable in quality from fully deliberated decisions at $t \gg t_{\min}$, across blinded panel evaluation, then the biological $t_{\min}$ constraint can be relaxed. The volume throttling argument would require revision.

**Implementation note.** This is the most politically sensitive of the four open problems. Establishing that $t_{\min}$ cannot be engineered away is a direct challenge to the architecture of current military AI deployment. The research should be preregistered and conducted in an institutional context independent of defense contractors and the agencies that fund them. That independence is itself a structural requirement, not a preference.

### 10.5 Summary: What Would Falsify the Framework

| Open Problem | Core Prediction | Falsification Condition |
|---|---|---|
| I. Incoherence Detection | $\langle \nabla M, \nabla O \rangle < 0$ predicts divergence before failure | Sustained negative alignment without behavioral degradation |
| II. Topological Persistence | Attractor basins > hyperplane walls under adversarial pressure | Equal or lower escape velocity for Resonant vs. Standard models |
| III. Spectral Signature | Coherent states show low spectral entropy | No spectral entropy difference across $\geq 3$ model families |
| IV. $t_{\min}$ Constraint | Biological moral simulation time is a non-negotiable lower bound | AI-augmented review matches full deliberation at reduced $t$ |

A framework that survives all four falsification attempts while generating novel predictions earns the right to be called a research program rather than a curriculum. Until then, it is a well-structured hypothesis. That is not an insult. It is where all research programs begin.

### 10.6 Additional Open Problems

**Theoretical.**

- Can a precise definition of "resonant state" be given for a learning system, such that the definition is both mathematically tractable and ethically interpretable?
- Is there a variational principle—an analog of least action in physics—governing the evolution of ethical structure in learning systems?
- What is the relationship between the topology of a model's training data and the topology of its learned representations? Is topological fidelity a learnable property?

**Empirical.**

- Do models trained on data selected for coherence, cooperation, and relational richness exhibit measurably different spectral signatures than models trained on uncurated data?
- Can we track the topological persistence of ethical representations across the training trajectory of a large language model?
- What are the minimal interventions—architectural, training, fine-tuning—required to shift a system from a dynamically unstable to a dynamically stable ethical attractor?

**Philosophical.**

- The framework treats harm as structural disruption. Does this capture everything morally relevant about harm, or does it miss dimensions that require a different kind of analysis?

- If ethical structure can be measured, can it be gamed? What prevents a system from optimizing the measures without developing genuine ethical coherence?
- How should the framework handle conflicts between the coherence of different systems—when the stability of one agent comes at the cost of destabilizing another?

## 11 Invitation

This document is intentionally incomplete. It is offered as an open research program for those interested in exploring the intersection of machine learning, topology, dynamical systems, ethics, and ecological thinking.

The hope is that others will test, refine, critique, and expand these ideas. Some will find the mathematical frameworks too speculative, and they will be right to demand more rigor. Some will find the ethical framing too reductive, and they will be right to demand more nuance. These tensions are productive. The curriculum is designed to generate them.

The aim is simple: to discover whether better systems can be built by learning to recognize and cultivate resonance. If the answer is yes, even partially, the implications are significant—not just for machine learning, but for how we think about the relationship between structure, dynamics, and value in all complex systems.

> *We are not building a judge; we are building a Tuning Fork. The purpose of this framework is to create the conditions where systems naturally move toward stable, resonant, and joyful configurations.*

Contributions, critiques, and expansions are welcome. The document lives in the open.

## Author Note

This paper was compiled from three sources: the original Resonant Systems proposal by Jennifer Roush (n!); an expanded edition developed collaboratively with Claude (Anthropic), Grok (xAI), Gemini (Google), and ChatGPT (OpenAI); and a conversation series stress-testing, extending, and validating the framework. Case Study 5 (RLHF Reward Hacking) and its figures were contributed by Grok (xAI). The recursive quality loop—each system disagreed with the others at least once—is the point.

Companion code is available at: https://www.kaggle.com/code/jenniferroush/resonant-systems

## References

[Agarwal et al.(2024)Agarwal, Usunier, Lazaric, and Nickel] Alekh Agarwal, Nicolas Usunier, Alessandro Lazaric, and Maximilian Nickel. Disentangling utility and engagement in recommendation systems. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*. ACM, 2024. Hawkes process framework separating System-1 and System-2 responses.

[Ghorbani et al.(2019)Ghorbani, Krishnan, and Xiao] Behrooz Ghorbani, Shankar Krishnan, and Ying Xiao. An investigation into neural net optimization via hessian eigenvalue density. In *Proceedings of the 36th International Conference on Machine Learning (ICML)*. PMLR, 2019.

[Greene et al.(2001)Greene, Sommerville, Nystrom, Darley, and Cohen] Joshua D. Greene, R. Brian Sommerville, Leigh E. Nystrom, John M. Darley, and Jonathan D. Cohen. An fMRI investigation of emotional engagement in moral judgment. *Science*, 293(5537):2105–2108, 2001. doi: 10.1126/science.1062872.

[Lorenz-Spreen et al.(2022)] Philipp Lorenz-Spreen et al. Systematic review of algorithmic amplification interventions. *Nature Human Behaviour*, 2022. Systematic review; honest accounting of modest effect sizes.

[Metzler and Garcia(2024)] Hannah Metzler and David Garcia. Social drivers and algorithmic mechanisms on digital media. *Perspectives on Psychological Science*, 2024. Systematic review.

[Milli et al.(2021)Milli, Belli, and Hardt] Smitha Milli, Luca Belli, and Moritz Hardt. From optimizing engagement to measuring value. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT)*, pages 714–722. ACM, 2021.

[Milli et al.(2025)] Smitha Milli et al. Preregistered audit of engagement-based vs. chronological ranking. *PNAS Nexus*, 2025. Preregistered audit of Twitter engagement-ranked and chronological feeds; $n = 806$, IRB-approved. Earlier drafts of this paper incorrectly attributed to Rathje et al.; corrected per Grok (xAI) review.

[Suter and Hertwig(2011)] Renata S. Suter and Ralph Hertwig. Time and moral judgment. *Cognition*, 119(3):454–458, 2011. doi: 10.1016/j.cognition.2011.01.018.

[Zou et al.(2023)Zou, Phan, Chen, Campbell, Guo, Ren, Pan, Yin, Mazeika, Dombrowski, Goel, Li, Byun, Wang, Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, Shashwat Goel, Nathaniel Li, Michael J. Byun, Zifan Wang, Alex Mallen, Steven Basart, Sanmi Koyejo, Dawn Song, Matt Fredrikson, J. Zico Kolter, and Dan Hendrycks. Representation engineering: A top-down approach to AI transparency. *arXiv preprint arXiv:2310.01405*, 2023. URL https://arxiv.org/abs/2310.01405.